


Funktion

Funktion

1. Introduction	2
2. Installing Funktion	3
3. Getting Started	4
4. Using the CLI	6
4.1. Browsing resources	6
4.2. Deleting resources	6
4.3. Installing Runtimes and Connectors	6
4.4. Configuring Connectors	7
4.5. Running the Operator	7
4.6. Subscribing to events	8
5. Using funktion on the JVM.....	9
5.1. Examples	9
5.2. Getting started with Funktion and the JVM	9
6. How it works	14
6.1. Kubernetes Resources.....	14
6.2. Debugging	14
6.3. Terminology	15
7. FAQ	17
7.1. General questions	17

Chapter 1. Introduction

Funktion is an open source event driven lambda style programming model designed for [Kubernetes](#).



Funktion supports over 200 of different [event sources and connectors](#) including most network protocols, transports, databases, messaging systems, social networks, cloud services and SaaS offerings.

In a sense funktion is a [serverless](#) approach to event driven microservices as you focus on just writing simple *functions* in whatever programming language you prefer, then **funktion** and Kubernetes takes care of the rest. Its not that there's no servers; its more that you as the funktion developer don't have to worry about managing them!

Chapter 2. Installing Funktion

To use **funktion** you will need a [kubernetes](#) or [openshift](#) cluster.

If you are on your laptop a quick way to get a kubernetes cluster is by [installing and starting minikube](#) and then [installing kubectl](#) and putting it on your **PATH** environment variable.

You will also need to [download the funktion binary for your platform](#) and add it to your **PATH** environment variable

Chapter 3. Getting Started

Type the following commands.

To make it easier to see what kubernetes resources are being created you may wish to create a new namespace for this experiment first:

```
kubectl create namespace funky  
kubectl config set-context `kubectl config current-context` --namespace=funky
```

Now we'll install the runtimes and a couple of connectors

```
funktion install timer twitter
```

Now lets run the **funktion operator** to watch for funktion resources and create the necessary kubernetes **Deployment** and **Services**.

```
funktion operate
```

Open another terminal then type:

```
kubectl apply -f https://raw.githubusercontent.com/fabric8io/funktion-  
operator/master/examples/subscription1.yml
```

You should now have created a subscription flow. You can view the subscription via

```
funktion get subscription
```

To view the output of the subscription you can use the following (assuming you've [enabled tab completion for kubectl](#))

```
kubectl logs -f subscription1-[TAB]
```

If you don't have tab completion you can specify the exact pod name, or you can use this command to find it and use it:

```
kubectl logs -f `kubectl get pod -oname -lfunktion.fabric8.io/kind=Subscription`
```

To delete the subscription:

```
funktion delete subscription subscription1
```

Now lets create a function:

```
kubectl apply -f https://raw.githubusercontent.com/fabric8io/funktion-  
operator/master/examples/function1.yml
```

If you are running the [fabric8 console](#) then you will have the link:[exposecontroller] microservice running and will be able to invoke it via running one of these commands:

```
minikube service funktion1 -n funky  
gofabric8 service funktion1 -n funky
```

Or clicking on the `funktion1` service in the [fabric8 console](#) in the `Services` tab for the `funky` namespace.

Chapter 4. Using the CLI

You can get help on the available commands via:

```
funktion
```

4.1. Browsing resources

To list all the resources of different kind via:

```
funktion get connector  
funktion get subscription  
funktion get function  
funktion get runtime
```

or to save typing you can use:

```
funktion get c  
funktion get s  
funktion get f  
funktion get r
```

4.2. Deleting resources

You can delete a Connector or Subscription via:

```
funktion delete connector foo  
funktion delete subscription bar  
funktion delete function whatnot  
funktion delete runtime nodejs
```

Or to remove all the Subscriptions or Connectors use `--all`

```
funktion delete subscription --all
```

4.3. Installing Runtimes and Connectors

To install the default function runtimes and connectors into your namespace type the following:

```
funktion install --all-connectors
```

There's over [200 connectors](#) provided out of the box. If you only want to install a number of them

you can specify their names as parameters

```
funktion install amqp kafka timer twitter
```

To just get a feel for what connectors are available without installing them try:

```
funktion install --list-connectors
```

or for short:

```
funktion install -l
```

4.4. Configuring Connectors

Various connectors have different configuration properties. For example the `twitter` connector has a number of properties to configure like the secret and token.

So to configure a connector you can type:

```
funktion edit connector twitter
```

You will then be prompted to enter new values; you can just hit `[ENTER]` to avoid changing a property.

To see a list of all the properties you can type

```
funktion edit connector twitter -l
```

Then you can pass in specific properties directly via the non-interactive version of the edit command:

```
funktion edit connector twitter accessToken=mytoken accessTokenSecret=mysecret  
consumerKey=myconsumerkey consumerSecret=myconsumerSecert
```

4.5. Running the Operator

You can run the funktion operator from the command line if you prefer:

```
funktion operate
```

Though ideally we'd run the `funktion application` inside kubernetes; via a helm chart, `kubect1 apply` or the `Run...` button in the [fabric8 developer console](#)

4.6. Subscribing to events

To create a new subscription for a connector try the following:

```
funktion subscribe --from timer://bar?period=5000 --to http://foo/
```

This will generate a new **Subscription** which will result in a new **Deployment** being created and one or more Pods should spin up.

Note that you must be running the **Operator** as described in the section above; its the **Operator** which actually creates a **Deployment** for each **Subscription**.

Also note that the first time you try out a new Connector kind it may take a few moments to download the docker image for this connector - particularly the first time you use a connector.

Once a pod has started for the **Deployment** you can then view the logs of a subscription via **kubectl**

```
kubectl logs -f nameOfSubscription[TAB]
```

Scaling a Subscription

If you want to stop a subscription type:

```
kubectl scale --replicas=0 deployment nameOfSubscription
```

To start it again:

```
kubectl scale --replicas=1 deployment nameOfSubscription
```

Using kubectl directly

You can also create a Subscription using **kubectl** if you prefer:

```
kubectl apply -f https://github.com/fabric8io/funktion-operator/blob/master/examples/subscription1.yml
```

You can view all the Connectors and Subscriptions via:

```
kubectl get cm
```

Or delete them via

```
kubectl delete cm nameOfConnectorOrSubscription
```

Chapter 5. Using funktion on the JVM

Funktion is designed so that it can bind any events to any HTTP endpoint or any function source using a scripting language like nodejs, python or ruby. But you can also embed the funktion mechanism inside a JVM process.

To do that you:

- write a simple function in any programming language [like this](#).
- create a [funktion.yml](#) file and associate your function with an [event trigger endpoint URL](#) such as a HTTP URL or email address to listen on, a message queue name or database table etc.
- build and deploy the Java project in the usual way, such as via [Jenkins CI / CD pipeline](#) and your funktion will be deployed to your kubernetes cluster!

5.1. Examples

Check out the following example projects which use a JVM and implement the functions in different JVM based languages:

- [funktion-java-example](#) is an example using a Java funktion triggered by HTTP
- [funktion-groovy-example](#) is an example using a [Groovy](#) funktion triggered by HTTP
- [funktion-kotlin-example](#) is an example using a [Kotlin](#) funktion triggered by HTTP

5.2. Getting started with Funktion and the JVM

You can just fork one of the above examples and use command line tools to build and deploy it to a [Kubernetes](#) or [OpenShift](#) cluster.

However to make it easier to create, build, test, stage, approve, release, manage and iterate on your funktion code from inside your browser we recommend you use the [Fabric8 Microservices Platform](#) with its baked in [Continuous Delivery](#) based on [Jenkins Pipelines](#) together with integrated [Developer Console](#), [Management](#) (centralised logging, metrics, alerts), [ChatOps](#) and [Chaos Monkey](#).

When using the [Fabric8 Microservices Platform](#) you can create a new funktion in a few clicks from the [Create Application](#) button; then the platform takes care of building, testing, staging and approving your releases, rolling upgrades, management and monitoring; you just use your browser via the [Developer Console](#) to create, edit or test your code while funktion, Jenkins and Kubernetes take care of building, packaging, deploying, testing and releasing your project.

5.2.1. Using the Fabric8 Microservices Platform

First you will need to install the [fabric8 microservices platform](#) on a cluster of [Kubernetes](#) (1.2 or later) or [OpenShift](#) (3.2 or later).

- follow one of the [fabric8 getting started guides](#) to get the [fabric8 microservices platform](#) up and running on a Kubernetes or OpenShift cluster

- open the [Developer Console](#)
- select your [Team Dashboard](#) page

5.2.2. Create and use your funktion

- from inside your [Team Dashboard](#) page click [Create Application](#) button then you will be presented with a number of different kinds of microservice to create
- select the [Funktion](#) icon and type in the name of your microservice and hit [Next](#)

- select the kind of funktion you wish to create (Java, Groovy, Kotlin, NodeJS etc) then hit **Next**
- you will now be prompted to choose one of the default CD Pipelines to use. For your first funktion we recommend **CanaryReleaseAndStage**
- selecting **Copy pipeline to project** is kinda handy if you want to edit your **Jenkinsfile** from your source code later on

- click **Next** then your app should be built and deployed. Please be patient first time you build a funktion as its going to be downloading a few docker images to do the build and runtime. You're second build should be much faster!
- once the build is complete you should see on the **App Dashboard** page the build pipeline run, the running pods for your funktion in each environment for your CD Pipeline and a link so you can easily navigate to the environment or ReplicaSet/ReplicationController/Pods in kubernetes
- in the screenshot below you can see we're running version **1.0.1** of the app **groovyfunktion** which currently has **1** running pod (those are all clickable links to view the ReplicationController or pods)
- for HTTP based funktions you can invoke the funktion via the open icon in the **Staging** environment (the icon to the right of the green **1** button next to **groovyfunktion-1: 1.0.1**)

5.2.3. How it works

When you implement your **Funktion** using a JVM based language like Java, Groovy, Kotlin or Scala then your function is packaged up into a [Spring Boot](#) application using [Apache Camel](#) to implement the trigger via the various [endpoint URLs](#).

We've focussed **funktion** on being some simple declarative metadata to describe triggers via URLs and a simple programming model which is the only thing funktion developers should focus on; leaving the implementation free to use different approaches for optimal resource usage.

The creation of the docker images and generation of the kubernetes manifests is all done by the [fabric8-maven-plugin](#) which can work with pure docker on Kubernetes or reuse OpenShift's binary source to image builds. Usually this is hidden from you if you are using the [Continuous Delivery](#) in the [fabric8 microservices platform](#); but if you want to play with funktion purely from the command line, you'll need to [install Java](#) and [install Apache Maven](#).

Underneath the covers a [Kubernetes Deployment](#) is automatically created for your Funktion (or on OpenShift a [DeploymentConfig](#) is used) which takes care of scaling your funktion and performing [rolling updates](#) as you edit your code.

Chapter 6. How it works

The `funktion operator` watches for `Subscription` and `Function` resources.

When a new `Subscription` is created then this operator will spin up a matching `Deployment` which consumes from some `Connector` and typically invokes a function using HTTP.

When a new is created then this operator will spin up a matching `Deployment` for running the `function source code` along with a `Service` to expose the service as a HTTP or HTTPS endpoint.

The following kubernetes resources are used:

6.1. Kubernetes Resources

A `Subscription` is modelled as a Kubernetes `ConfigMap` with the label `kind.funktion.fabric8.io: "Subscription"`. A `ConfigMap` is used so that the entries inside the `ConfigMap` can be mounted as files inside the `Deployment`. For example this will typically involve storing the `funktion.yml` file or maybe a Spring Boot `application.properties` file inside the `ConfigMap` like [this example subscription](#)

A `Connector` is generated [for every Camel Component](#) and each connector has an associated `ConfigMap` resource like [this example](#) which uses the label `kind.funktion.fabric8.io: "Connector"`. The `Connector` stores the `Deployment` metadata, the `schema.yml` for editing the connectors endpoint URL and the `documentation.adoc` documentation for using the Connector.

So a `Connector` can have `0..N Subscriptions` associated with it. For those who know [Apache Camel](#) this is like the relationship between a `Component` having `0..N Endpoints`.

For example we could have a Connector called `kafka` which knows how to produce and consume messages on [Apache Kafka](#) with the Connector containing the metadata of how to create a consumer, how to configure the kafka endpoint and the documentnation. Then a Subscription could be created for `kafka://cheese` to subscribe on the `cheese` topic and post messages to [http://foo/](#).

Typically a number of `Connector` resources are shipped as a package; such as inside the [Red Hat iPaaS](#) or as an app inside fabric8. Though a `Connector` can be created as part of the CD Pipeline by an expert Java developer who takes a Camel component and customizes it for use by `Funktion` or the `iPaaS`.

The collection of `Connector` resources installed in a kubernetes namespace creates the `integration palette` thats seen by users in tools like CLI or web UIs.

Then a `Subscription` can be created at any time by users from a `Connector` with a custom configuration (e.g. choosing a particular queue or topic in a messaging system or a particular table in a database or folder in a file system).

6.2. Debugging

If you ever need to you can debug any `Subscription` as each Subscription matches a `Deployment` of one or more pods. So you can just debug that pod which typically is a regular Spring Boot and camel application.

Otherwise you can debug the pod that's exposing an HTTP endpoint using whatever the native debugger is; e.g. using Java or NodeJS or whatever.

6.3. Terminology

This section defines all the terms used in the `funktion` project

6.3.1. Connector

A `connector` represents a way to connect to some event source, including most network protocols, transports, databases, messaging systems, social networks, cloud services and SaaS offerings. Funktion supports [over 200 event sources](#).

At the implementation level a `connector` represents the kubernetes `deployment` metadata required to take the `flow` and implement it as one or more kubernetes `pods`.

6.3.2. Flow

A `flow` is a sequence of `steps` such as consuming events from an `endpoint` or invoking a `function`.

6.3.3. Subscription

A subscription consists of one or more `flows` which bind events to HTTP endpoints and functions.

For example here is a sample subscription with a single flow:

```
flows:
- steps:
  - kind: endpoint
    uri: timer://foo?fixedRate=true&period=5000
  - kind: endpoint
    uri: http://myendpoint/
```

Creating a `subscription` results in the `funktion operator` creating an associated `deployment` which implements the flows.

6.3.4. Function

A `function` is some source code to implement a function in some programming language like JavaScript, python or ruby.

6.3.5. Runtime

A `runtime` represents the kubernetes `deployment` metadata required to take a function source in some programming language and implement it as one or more pods.

The `funktion operator` then detects a new `function` resource being created or updated and creates the associated `runtime` deployment

6.3.6. Funktion Operator

The `funktion operator` is a running `pod` in kubernetes which monitors for all the funktion resources like `function`, `runtime`, `connector` and `subscription` and creates, updates or deletes the associated kubernetes `deployments` and `services` so that as you create a `subscription` or `function` the associated kubernetes resources are created.

Chapter 7. FAQ

Here are the frequently asked questions:

7.1. General questions

7.1.1. What is the license?

All of the Funktion source code is licensed under the [Apache License 2.0](<https://www.apache.org/licenses/LICENSE-2.0>)

7.1.2. How do I get started?

Please [Install Funktion](#) then follow the [Getting Started Guide](#)

7.1.3. How do I install funktion?

See the [Install Guide](#)

7.1.4. What is serverless?

The term **serverless** just means that with lambda style programming the developer just focuses on writing **functions** only - there is no need for developers to think about managing servers or even containers.

Its not that there are no servers - of course there are - its just that developers don't need to think about them at all, they are managed for you by the platform.