

Funktion


# Funktion

1. Introduction .....	2
2. Installing Funktion .....	3
3. Getting Started .....	4
3.1. A more complex example.....	5
4. Using the CLI .....	6
4.1. Browsing resources .....	6
4.2. Deleting resources .....	6
4.3. Installing Runtimes and Connectors .....	6
4.4. Configuring Connectors .....	7
4.5. Running the Operator .....	7
4.6. Subscribing to events .....	8
5. Using funktion on the JVM.....	9
5.1. Examples .....	9
5.2. Getting started with Funktion and the JVM.....	9
6. How it works.....	14
6.1. Kubernetes Resources.....	14
6.2. Debugging .....	14
6.3. Terminology .....	15
7. FAQ .....	17
7.1. General questions .....	17



# Chapter 1. Introduction

**Funktion** is an open source event driven lambda style programming model designed for [Kubernetes](#).



Funktion supports over 200 of different [event sources and connectors](#) including most network protocols, transports, databases, messaging systems, social networks, cloud services and SaaS offerings.

In a sense funktion is a [serverless](#) approach to event driven microservices as you focus on just writing simple *functions* in whatever programming language you prefer, then **funktion** and Kubernetes takes care of the rest. Its not that there's no servers; its more that you as the funktion developer don't have to worry about managing them!

# Chapter 2. Installing Funktion

To use **funktion** you will need a **kubernetes** or **openshift** cluster.

If you are on your laptop a quick way to get a kubernetes cluster is by **installing and starting minikube** and then **installing kubectl** and putting it on your **PATH** environment variable.

You will also need to **download the funktion binary for your platform** and add it to your **PATH** environment variable

# Chapter 3. Getting Started

Type the following commands.

To make it easier to see what kubernetes resources are being created you may wish to create a new namespace for this experiment first:

```
kubectl create namespace funky  
kubectl config set-context `kubectl config current-context` --namespace=funky
```

Now we'll install the runtimes and a couple of connectors

```
funktion install http4 timer twitter
```

Now lets run the **funktion operator** to watch for funktion resources and create the necessary kubernetes **Deployment** and **Services**.

```
funktion operate
```

Open another terminal then type:

```
kubectl apply -f https://raw.githubusercontent.com/fabric8io/funktion-  
operator/master/examples/subscription1.yml
```

You should now have created a subscription flow. You can view the subscription via

```
funktion get subscription
```

To view the output of the subscription you can use the following (assuming you've [enabled tab completion for kubectl](#))

```
kubectl logs -f subscription1-[TAB]
```

If you don't have tab completion you can specify the exact pod name, or you can use this command to find it and use it:

```
kubectl logs -f `kubectl get pod -oname -lfunktion.fabric8.io/kind=Subscription`
```

To delete the subscription:

```
funktion delete subscription subscription1
```

Now lets create a function:

```
kubectl apply -f https://raw.githubusercontent.com/fabric8io/funktion-  
operator/master/examples/function1.yml
```

If you are running the [fabric8 console](#) then you will have the link:[exposecontroller] microservice running and will be able to invoke it via running one of these commands:

```
minikube service funktion1 -n funky  
gofabric8 service funktion1 -n funky
```

Or clicking on the `funktion1` service in the [fabric8 console](#) in the `Services` tab for the `funky` namespace.

### 3.1. A more complex example

To see a more real world style example check out the [blog splitting and counting example using functions and flows](#)

# Chapter 4. Using the CLI

You can get help on the available commands via:

```
funktion
```

## 4.1. Browsing resources

To list all the resources of different kind via:

```
funktion get connector  
funktion get subscription  
funktion get function  
funktion get runtime
```

or to save typing you can use:

```
funktion get c  
funktion get s  
funktion get f  
funktion get r
```

## 4.2. Deleting resources

You can delete a Connector or Subscription via:

```
funktion delete connector foo  
funktion delete subscription bar  
funktion delete function whatnot  
funktion delete runtime nodejs
```

Or to remove all the Subscriptions or Connectors use `--all`

```
funktion delete subscription --all
```

## 4.3. Installing Runtimes and Connectors

To install the default function runtimes and connectors into your namespace type the following:

```
funktion install --all-connectors
```

There's over [200 connectors](#) provided out of the box. If you only want to install a number of them



you can specify their names as parameters

```
funktion install amqp kafka timer twitter
```

To just get a feel for what connectors are available without installing them try:

```
funktion install --list-connectors
```

or for short:

```
funktion install -l
```

## 4.4. Configuring Connectors

Various connectors have different configuration properties. For example the `twitter` connector has a number of properties to configure like the secret and token.

So to configure a connector you can type:

```
funktion edit connector twitter
```

You will then be prompted to enter new values; you can just hit `[ENTER]` to avoid changing a property.

To see a list of all the properties you can type

```
funktion edit connector twitter -l
```

Then you can pass in specific properties directly via the non-interactive version of the edit command:

```
funktion edit connector twitter accessToken=mytoken accessTokenSecret=mysecret  
consumerKey=myconsumerkey consumerSecret=myconsumerSecert
```

## 4.5. Running the Operator

You can run the funktion operator from the command line if you prefer:

```
funktion operate
```

Though ideally we'd run the `funktion application` inside kubernetes; via a helm chart, `kubect1 apply` or the `Run...` button in the [fabric8 developer console](#)

## 4.6. Subscribing to events

To create a new subscription for a connector try the following:

```
funktion subscribe timer://bar?period=5000 http://foo/
```

This will generate a new **Subscription** which will result in a new **Deployment** being created and one or more Pods should spin up.

Note that you must be running the **Operator** as described in the section above; its the **Operator** which actually creates a **Deployment** for each **Subscription**.

Also note that the first time you try out a new Connector kind it may take a few moments to download the docker image for this connector - particularly the first time you use a connector.

Once a pod has started for the **Deployment** you can then view the logs of a subscription via **kubectl**

```
kubectl logs -f nameOfSubscription[TAB]
```

### Scaling a Subscription

If you want to stop a subscription type:

```
kubectl scale --replicas=0 deployment nameOfSubscription
```

To start it again:

```
kubectl scale --replicas=1 deployment nameOfSubscription
```

### Using kubectl directly

You can also create a Subscription using **kubectl** if you prefer:

```
kubectl apply -f https://github.com/fabric8io/funktion-operator/blob/master/examples/subscription1.yml
```

You can view all the Connectors and Subscriptions via:

```
kubectl get cm
```

Or delete them via

```
kubectl delete cm nameOfConnectorOrSubscription
```

# Chapter 5. Using funktion on the JVM

Funktion is designed so that it can bind any events to any HTTP endpoint or any function source using a scripting language like nodejs, python or ruby. But you can also embed the funktion mechanism inside a JVM process.

To do that you:

- write a simple function in any programming language [like this](#).
- create a [funktion.yml](#) file and associate your function with an [event trigger endpoint URL](#) such as a HTTP URL or email address to listen on, a message queue name or database table etc.
- build and deploy the Java project in the usual way, such as via [Jenkins CI / CD pipeline](#) and your funktion will be deployed to your kubernetes cluster!

## 5.1. Examples

Check out the following example projects which use a JVM and implement the functions in different JVM based languages:

- [funktion-java-example](#) is an example using a Java funktion triggered by HTTP
- [funktion-groovy-example](#) is an example using a [Groovy](#) funktion triggered by HTTP
- [funktion-kotlin-example](#) is an example using a [Kotlin](#) funktion triggered by HTTP

## 5.2. Getting started with Funktion and the JVM

You can just fork one of the above examples and use command line tools to build and deploy it to a [Kubernetes](#) or [OpenShift](#) cluster.

However to make it easier to create, build, test, stage, approve, release, manage and iterate on your funktion code from inside your browser we recommend you use the [Fabric8 Microservices Platform](#) with its baked in [Continuous Delivery](#) based on [Jenkins Pipelines](#) together with integrated [Developer Console](#), [Management](#) (centralised logging, metrics, alerts), [ChatOps](#) and [Chaos Monkey](#).

When using the [Fabric8 Microservices Platform](#) you can create a new funktion in a few clicks from the [Create Application](#) button; then the platform takes care of building, testing, staging and approving your releases, rolling upgrades, management and monitoring; you just use your browser via the [Developer Console](#) to create, edit or test your code while funktion, Jenkins and Kubernetes take care of building, packaging, deploying, testing and releasing your project.

### 5.2.1. Using the Fabric8 Microservices Platform

First you will need to install the [fabric8 microservices platform](#) on a cluster of [Kubernetes](#) (1.2 or later) or [OpenShift](#) (3.2 or later).

- follow one of the [fabric8 getting started guides](#) to get the [fabric8 microservices platform](#) up and running on a Kubernetes or OpenShift cluster

- open the [Developer Console](#)
- select your [Team Dashboard](#) page

### 5.2.2. Create and use your funktion

- from inside your [Team Dashboard](#) page click [Create Application](#) button then you will be presented with a number of different kinds of microservice to create
- select the [Funktion](#) icon and type in the name of your microservice and hit [Next](#)

- select the kind of funktion you wish to create (Java, Groovy, Kotlin, NodeJS etc) then hit **Next**
- you will now be prompted to choose one of the default CD Pipelines to use. For your first funktion we recommend **CanaryReleaseAndStage**
- selecting **Copy pipeline to project** is kinda handy if you want to edit your **Jenkinsfile** from your source code later on

- click **Next** then your app should be built and deployed. Please be patient first time you build a funktion as its going to be downloading a few docker images to do the build and runtime. You're second build should be much faster!
- once the build is complete you should see on the **App Dashboard** page the build pipeline run, the running pods for your funktion in each environment for your CD Pipeline and a link so you can easily navigate to the environment or ReplicaSet/ReplicationController/Pods in kubernetes
- in the screenshot below you can see we're running version **1.0.1** of the app **groovyfunktion** which currently has **1** running pod (those are all clickable links to view the ReplicationController or pods)
- for HTTP based funktions you can invoke the funktion via the open icon in the **Staging** environment (the icon to the right of the green **1** button next to **groovyfunktion-1: 1.0.1**)

### 5.2.3. How it works

When you implement your **Funktion** using a JVM based language like Java, Groovy, Kotlin or Scala then your function is packaged up into a [Spring Boot](#) application using [Apache Camel](#) to implement the trigger via the various [endpoint URLs](#).

We've focussed **funktion** on being some simple declarative metadata to describe triggers via URLs and a simple programming model which is the only thing funktion developers should focus on; leaving the implementation free to use different approaches for optimal resource usage.

The creation of the docker images and generation of the kubernetes manifests is all done by the [fabric8-maven-plugin](#) which can work with pure docker on Kubernetes or reuse OpenShift's binary source to image builds. Usually this is hidden from you if you are using the [Continuous Delivery](#) in the [fabric8 microservices platform](#); but if you want to play with funktion purely from the command line, you'll need to [install Java](#) and [install Apache Maven](#).

Underneath the covers a [Kubernetes Deployment](#) is automatically created for your Funktion (or on OpenShift a [DeploymentConfig](#) is used) which takes care of scaling your funktion and performing [rolling updates](#) as you edit your code.

# Chapter 6. How it works

The `funktion operator` watches for `Subscription` and `Function` resources.

When a new `Subscription` is created then this operator will spin up a matching `Deployment` which consumes from some `Connector` and typically invokes a function using HTTP.

When a new is created then this operator will spin up a matching `Deployment` for running the `function source code` along with a `Service` to expose the service as a HTTP or HTTPS endpoint.

The following kubernetes resources are used:

## 6.1. Kubernetes Resources

A `Subscription` is modelled as a Kubernetes `ConfigMap` with the label `kind.funktion.fabric8.io: "Subscription"`. A `ConfigMap` is used so that the entries inside the `ConfigMap` can be mounted as files inside the `Deployment`. For example this will typically involve storing the `funktion.yml` file or maybe a Spring Boot `application.properties` file inside the `ConfigMap` like [this example subscription](#)

A `Connector` is generated [for every Camel Component](#) and each connector has an associated `ConfigMap` resource like [this example](#) which uses the label `kind.funktion.fabric8.io: "Connector"`. The `Connector` stores the `Deployment` metadata, the `schema.yml` for editing the connectors endpoint URL and the `documentation.adoc` documentation for using the Connector.

So a `Connector` can have `0..N` `Subscriptions` associated with it. For those who know [Apache Camel](#) this is like the relationship between a `Component` having `0..N` `Endpoints`.

For example we could have a Connector called `kafka` which knows how to produce and consume messages on [Apache Kafka](#) with the Connector containing the metadata of how to create a consumer, how to configure the kafka endpoint and the documetnation. Then a Subscription could be created for `kafka://cheese` to subscribe on the `cheese` topic and post messages to [http://foo/](#).

Typically a number of `Connector` resources are shipped as a package; such as inside the [Red Hat iPaaS](#) or as an app inside fabric8. Though a `Connector` can be created as part of the CD Pipeline by an expert Java developer who takes a Camel component and customizes it for use by `Funktion` or the `iPaaS`.

The collection of `Connector` resources installed in a kubernetes namespace creates the `integration palette` thats seen by users in tools like CLI or web UIs.

Then a `Subscription` can be created at any time by users from a `Connector` with a custom configuration (e.g. choosing a particular queue or topic in a messaging system or a particular table in a database or folder in a file system).

## 6.2. Debugging

If you ever need to you can debug any `Subscription` as each Subscription matches a `Deployment` of one or more pods. So you can just debug that pod which typically is a regular Spring Boot and camel application.



Otherwise you can debug the pod that's exposing an HTTP endpoint using whatever the native debugger is; e.g. using Java or NodeJS or whatever.

## 6.3. Terminology

This section defines all the terms used in the `funktion` project

### 6.3.1. Connector

A `connector` represents a way to connect to some event source, including most network protocols, transports, databases, messaging systems, social networks, cloud services and SaaS offerings. Funktion supports [over 200 event sources](#).

At the implementation level a `connector` represents the kubernetes `deployment` metadata required to take the `flow` and implement it as one or more kubernetes `pods`.

### 6.3.2. Flow

A `flow` is a sequence of `steps` such as consuming events from an `endpoint` or invoking a `function`.

### 6.3.3. Subscription

A subscription consists of one or more `flows` which bind events to HTTP endpoints and functions.

For example here is a sample subscription with a single flow:

```
flows:
- steps:
  - kind: endpoint
    uri: timer://foo?fixedRate=true&period=5000
  - kind: endpoint
    uri: http://myendpoint/
```

Creating a `subscription` results in the `funktion operator` creating an associated `deployment` which implements the flows.

### 6.3.4. Function

A `function` is some source code to implement a function in some programming language like JavaScript, python or ruby.

### 6.3.5. Runtime

A `runtime` represents the kubernetes `deployment` metadata required to take a function source in some programming language and implement it as one or more pods.

The `funktion operator` then detects a new `function` resource being created or updated and creates the associated `runtime` deployment

### 6.3.6. Funktion Operator

The `funktion operator` is a running `pod` in kubernetes which monitors for all the funktion resources like `function`, `runtime`, `connector` and `subscription` and creates, updates or deletes the associated kubernetes `deployments` and `services` so that as you create a `subscription` or `function` the associated kubernetes resources are created.

# Chapter 7. FAQ

Here are the frequently asked questions:

## 7.1. General questions

### 7.1.1. What is the license?

All of the Funktion source code is licensed under the [Apache License 2.0](<https://www.apache.org/licenses/LICENSE-2.0>)

### 7.1.2. How do I get started?

Please [Install Funktion](#) then follow the [Getting Started Guide](#)

### 7.1.3. How do I install funktion?

See the [Install Guide](#)

### 7.1.4. What is serverless?

The term **serverless** just means that with lambda style programming the developer just focuses on writing **functions** only - there is no need for developers to think about managing servers or even containers.

Its not that there are no servers - of course there are - its just that developers don't need to think about them at all, they are managed for you by the platform.